

HW4

Myungjin Lee

5128876730


1a)

Process :


1. Find the first leaf where $\text{age} \geq 10$ and $\text{age} \neq 15$ in root (in this case)
2. Go to the block which has keys of $\text{age} \geq 10$ and traverse next block through leaf node.
3. When the keys are satisfied with the range $\text{age} \geq 10$ then read the corresponding internal node except $\text{age} = 15$.

The number of blocks : 5

1b)



1c)



2-a) Nested-loop join with R as the outer relation

Step1 : Read R once (cost : $B(R)$)

Memory : $(M-2)$ blocks of R as a hash table(since 1 for input buffer for S and 1 for output buffer)

Step2 : The size of each run is $(M-2)$, so that the number of runs is $\lceil B(R)/(M-2) \rceil$.

Each time of $\lceil B(R)/(M-2) \rceil$ reads S (cost : $B(S) \lceil B(R)/(M-2) \rceil$)

=> $\lceil 500/98 \rceil = 6$ runs for R, each time for 6 runs reads S

Total cost: $B(R) + B(S) \lceil B(R)/(M-2) \rceil$

$$= 500 + 10,000 * \lceil 500/98 \rceil = 60,500$$

2-b) Nested-loop join with S as the outer relation

Step1 : Read S once (cost : $B(S)$)

Memory : $(M-2)$ blocks of S as a hash table(since 1 for input buffer for R and 1 for output buffer)

Step2 : The size of each run is $(M-2)$, so that the number of runs are $\lceil B(S)/(M-2) \rceil$.

Each time of $\lceil B(S)/(M-2) \rceil$ reads R (cost : $B(R) \lceil B(S)/(M-2) \rceil$)

=> $\lceil 10,000/98 \rceil = 103$ runs for S, each time for 103 runs reads R

Total cost: $B(S) + B(R) \lceil B(S)/(M-2) \rceil$

$$= 10,000 + 500 * \lceil 10,000/98 \rceil = 61,500$$

2-c) Sort-merge join

Step1 : Split R into runs of size M and split S into runs of size M. (cost : $2B(R) + 2B(S)$)

=> $500/100 = 5$ runs of size 100 for R, $10,000/100 = 100$ runs of size 100 for S

Step2 : Merge $M-1$ runs into a new run (cost : $2B(R) + 2B(S)$)

=> $\lceil 105/99 \rceil = 2$ runs

Step3 : Join 2 runs from R and S (cost : $B(R) + B(S)$)

Total cost: $2B(R) + 2B(S) + 2B(R) + 2B(S) + B(R) + B(S) = 5B(R) + 5B(S)$

$$= 5 * 500 + 5 * 10,000 = 52,500$$

2-d) Simple sort-based join

Sort R:

Step1 : Split R into runs of size M so that we get $B(R)/M$ lists : $500/100 = 5$ lists of size 100

Step2 : Merge 5 runs : 1 sorted list of size 500 (cost : $2B(R) * (\# \text{ of pass})$)

Sort S:

Step1 : Split S into runs of size M so that we get $B(S)/M$ lists

=> $10,000/100 = 100$ lists of size 100

Step2 : Merge M-1 runs into a new run (since 1 is for output buffer)

=> 1 sorted list of size 9,900 and 1 sorted list of size 100

Step3 : Merge two lists (1 size of 9,900 and 1 size of 100) (cost : $2B(S)*(\# \text{ of pass})$)

=> 1 sorted list size of 10,000

Finally, read both relations in sorted order, match tuples (join R and S) (cost : $B(R) + B(S)$)

Total cost: $2B(R)*(\# \text{ of pass}) + 2B(S)*(\# \text{ of pass}) + B(R) + B(S)$

= $2*500*2 + 2*10,000*3 + 500 + 10,000 = 72,500$

2-e) Partitioned-hash join

Step1 : Hash R and S into M-1 buckets. (since 1 buffer is for input) (cost : $2B(R) + 2B(S)$)

(# buckets : M-1, size of buckets : $B(R)/(M-1)$ for R and $B(S)/(M-1)$ for S)

=> 99 buckets of size 500/99 for R and 99 buckets of size 10,000/99 for S

Step2 : Load the buckets of R_i with corresponding S_i , one block at a time, output joining tuples with using another hash function

(cost : $B(R) + B(S)$)

Total cost: $2B(R) + 2B(S) + B(R) + B(S) = 3B(R) + 3B(S)$

= $3*500 + 3*10,000 = 31,500$

2-f) Index join

Step1 : Iterate over R, for each tuple, fetch corresponding tuples from S

(# joined tuples : $B(S)/V(S,a)$ from S)

Total cost: $B(R) + B(S)T(R)/V(S,a)$

= $500 + 10,000*5,000/200 = 250,500$

The most efficient algorithm : Partitioned-hash join